

# Assignment-05-Multiple Linear Regression-1

**Q1.Consider only the below columns and prepare a prediction model for predicting Price.**

```
Corolla<-Corolla[c("Price","Age_08_04","KM","HP","cc","Doors","Gears","Quarterly_Tax","Weight")]
```

In [1]:

```
# Import Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.formula.api as smf
import statsmodels.api as sm
from statsmodels.graphics.regressionplots import influence_plot
```

In [2]:

```
# Import dataset
toyo=pd.read_csv("C:/Users/LENOVO/Documents/Custom Office Templates/ToyotaCorolla.csv",enco
toyo
```

Out[2]:

	</										

	Id	Model	Price	Age_08_04	Mfg_Month	Mfg_Year	KM	Fuel_Type	HP	Met_
1433	1440	TOYOTA	8500	71	10	1998	17016	Petrol	86	
		Corolla								
		1.3 16V								
1434	1441	HATCHB	7250	70	11	1998	16916	Petrol	86	
		LINEA								
		TERRA								
1435	1442	TOYOTA	6950	76	5	1998	1	Petrol	110	
		Corolla								
		1.6 LB								
		LINEA								
		TERRA								
		4/5-								
		Doors								

1436 rows × 38 columns

EDA

In [3]:

```
toyo.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 38 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Id                    1436 non-null   int64
 1   Model                 1436 non-null   object
 2   Price                 1436 non-null   int64
 3   Age_08_04            1436 non-null   int64
 4   Mfg_Month            1436 non-null   int64
 5   Mfg_Year             1436 non-null   int64
 6   KM                   1436 non-null   int64
 7   Fuel_Type            1436 non-null   object
 8   HP                   1436 non-null   int64
 9   Met_Color            1436 non-null   int64
10   Color                1436 non-null   object
11   Automatic            1436 non-null   int64
12   cc                   1436 non-null   int64
13   Doors                1436 non-null   int64
14   Cylinders            1436 non-null   int64
15   Gears                1436 non-null   int64
16   Quarterly_Tax        1436 non-null   int64
17   Weight               1436 non-null   int64
18   Mfr_Guarantee         1436 non-null   int64
19   BOVAG_Guarantee       1436 non-null   int64
20   Guarantee_Period      1436 non-null   int64
21   ABS                  1436 non-null   int64
22   Airbag_1             1436 non-null   int64
23   Airbag_2             1436 non-null   int64
24   Airco                1436 non-null   int64
25   Automatic_airco       1436 non-null   int64
26   Boardcomputer         1436 non-null   int64
27   CD_Player            1436 non-null   int64
28   Central_Lock         1436 non-null   int64
29   Powered_Windows      1436 non-null   int64
30   Power_Steering       1436 non-null   int64
31   Radio                1436 non-null   int64
32   Mistlamps            1436 non-null   int64
33   Sport_Model          1436 non-null   int64
34   Backseat_Divider     1436 non-null   int64
35   Metallic_Rim         1436 non-null   int64
36   Radio_cassette       1436 non-null   int64
37   Tow_Bar              1436 non-null   int64
dtypes: int64(35), object(3)
memory usage: 426.4+ KB
```

In [4]:

```
toyo.isnull().sum()
```

Out[4]:

Id	0
Model	0
Price	0
Age_08_04	0
Mfg_Month	0
Mfg_Year	0
KM	0
Fuel_Type	0
HP	0
Met_Color	0
Color	0
Automatic	0
cc	0
Doors	0
Cylinders	0
Gears	0
Quarterly_Tax	0
Weight	0
Mfr_Guarantee	0
BOVAG_Guarantee	0
Guarantee_Period	0
ABS	0
Airbag_1	0
Airbag_2	0
Airco	0
Automatic_airco	0
Boardcomputer	0
CD_Player	0
Central_Lock	0
Powered_Windows	0
Power_Steering	0
Radio	0
Mistlamps	0
Sport_Model	0
Backseat_Divider	0
Metallic_Rim	0
Radio_cassette	0
Tow_Bar	0

dtype: int64

Inference:

No NA values

In [5]:

```
toyo2=pd.concat([toyo.iloc[:,2:4],toyo.iloc[:,6:7],toyo.iloc[:,8:9],toyo.iloc[:,12:14],toyo
toyo2
```

Out[5]:

	Price	Age_08_04	KM	HP	cc	Doors	Gears	Quarterly_Tax	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...	...	...	...	...	...	...	...	...	...
1431	7500	69	20544	86	1300	3	5	69	1025
1432	10845	72	19000	86	1300	3	5	69	1015
1433	8500	71	17016	86	1300	3	5	69	1015
1434	7250	70	16916	86	1300	3	5	69	1015
1435	6950	76	1	110	1600	5	5	19	1114

1436 rows × 9 columns

In [6]:

```
# rename columns
toyo3=toyo2.rename({'Age_08_04':'Age','cc':'CC','Quarterly_Tax':'QT'},axis=1)
toyo3
```

Out[6]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...	...	...	...	...	...	...	...	...	...
1431	7500	69	20544	86	1300	3	5	69	1025
1432	10845	72	19000	86	1300	3	5	69	1015
1433	8500	71	17016	86	1300	3	5	69	1015
1434	7250	70	16916	86	1300	3	5	69	1015
1435	6950	76	1	110	1600	5	5	19	1114

1436 rows × 9 columns

In [7]:

```
toyo3[toyo3.duplicated()]
```

Out[7]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
113	24950	8	13253	116	2000	5	5	234	1320

In [8]:

```
toyo4=toyo3.drop_duplicates().reset_index(drop=True)  
toyo4
```

Out[8]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...	...	...	...	...	...	...	...	...	...
1430	7500	69	20544	86	1300	3	5	69	1025
1431	10845	72	19000	86	1300	3	5	69	1015
1432	8500	71	17016	86	1300	3	5	69	1015
1433	7250	70	16916	86	1300	3	5	69	1015
1434	6950	76	1	110	1600	5	5	19	1114

1435 rows × 9 columns

In [9]:

```
toyo4.describe()
```

Out[9]:

	Price	Age	KM	HP	CC	Doors	
count	1435.000000	1435.000000	1435.000000	1435.000000	1435.000000	1435.000000	1435
mean	10720.915679	55.980488	68571.782578	101.491986	1576.560976	4.032753	4.03
std	3608.732978	18.563312	37491.094553	14.981408	424.387533	0.952667	0.95
min	4350.000000	1.000000	1.000000	69.000000	1300.000000	2.000000	2.00
25%	8450.000000	44.000000	43000.000000	90.000000	1400.000000	3.000000	3.00
50%	9900.000000	61.000000	63451.000000	110.000000	1600.000000	4.000000	4.00
75%	11950.000000	70.000000	87041.500000	110.000000	1600.000000	5.000000	5.00
max	32500.000000	80.000000	243000.000000	192.000000	16000.000000	5.000000	5.00

## Correlation Analysis

In [10]:

```
toyo4.corr()
```

Out[10]:

	Price	Age	KM	HP	CC	Doors	Gears	QT
Price	1.000000	-0.876273	-0.569420	0.314134	0.124375	0.183604	0.063831	0.211508
Age	-0.876273	1.000000	0.504575	-0.155293	-0.096549	-0.146929	-0.005629	-0.193319
KM	-0.569420	0.504575	1.000000	-0.332904	0.103822	-0.035193	0.014890	0.283312
HP	0.314134	-0.155293	-0.332904	1.000000	0.035207	0.091803	0.209642	-0.302287
CC	0.124375	-0.096549	0.103822	0.035207	1.000000	0.079254	0.014732	0.305982
Doors	0.183604	-0.146929	-0.035193	0.091803	0.079254	1.000000	-0.160101	0.107353
Gears	0.063831	-0.005629	0.014890	0.209642	0.014732	-0.160101	1.000000	-0.005125
QT	0.211508	-0.193319	0.283312	-0.302287	0.305982	0.107353	-0.005125	1.000000
Weight	0.575869	-0.466484	-0.023969	0.087143	0.335077	0.301734	0.021238	0.621988




In [11]:

```
sns.set_style(style='darkgrid')  
sns.pairplot(toyo4)
```

Out[11]:

&lt;seaborn.axisgrid.PairGrid at 0x74dd8e9760&gt;



## Model Building

In [12]:

```
model=smf.ols("Price~Age+KM+HP+CC+Doors+Gears+QT+Weight",data=toyo4).fit()
```

In [13]:

```
# Finding coefficient parameters  
model.params
```

Out[13]:

```
Intercept    -5472.540368  
Age           -121.713891  
KM            -0.020737  
HP            31.584612  
CC            -0.118558  
Doors         -0.920189  
Gears         597.715894  
QT            3.858805  
Weight        16.855470  
dtype: float64
```

In [14]:

```
# Finding tvalues and pvalues  
model.tvalues , np.round(model.pvalues,5)
```

Out[14]:

```
(Intercept    -3.875273  
Age           -46.551876  
KM            -16.552424  
HP            11.209719  
CC            -1.316436  
Doors         -0.023012  
Gears         3.034563  
QT            2.944198  
Weight        15.760663  
dtype: float64,  
Intercept     0.00011  
Age           0.00000  
KM            0.00000  
HP            0.00000  
CC            0.18824  
Doors         0.98164  
Gears         0.00245  
QT            0.00329  
Weight        0.00000  
dtype: float64)
```

Inference:

CC and Doors are insignificant

In [15]:

```
# Finding rsquared values
model.rsquared ,model.rsquared_adj
```

Out[15]:

```
(0.8625200256947, 0.8617487495415146)
```

Inference:

R\_Squared value is 0.86 model is a Good model

In [16]:

```
# Build SLR and MLR for insignificant variables 'CC' and 'Doors'
# Also find their tvalues and pvalues
```

In [17]:

```
slr_c=smf.ols("Price~CC",data=toyo4).fit()
slr_c.tvalues, slr_c.pvalues # CC has significant pvalue
```

Out[17]:

```
(Intercept    24.879592
CC             4.745039
dtype: float64,
Intercept    7.236022e-114
CC           2.292856e-06
dtype: float64)
```

In [18]:

```
slr_d=smf.ols("Price~Doors",data=toyo4).fit()
slr_d.tvalues, slr_d.pvalues # Doors has significant value
```

Out[18]:

```
(Intercept    19.421546
Doors         7.070520
dtype: float64,
Intercept    8.976407e-75
Doors        2.404166e-12
dtype: float64)
```

In [19]:

```
mlr_cd=smf.ols("Price~CC+Doors",data=toyo4).fit()
mlr_cd.tvalues, mlr_cd.pvalues # CC and Doors have significant pvalue
```

Out[19]:

```
(Intercept    12.786341
CC             4.268006
Doors         6.752236
dtype: float64,
Intercept    1.580945e-35
CC           2.101878e-05
Doors        2.109558e-11
dtype: float64)
```

# Model Validation Techniques

## Two Techniques: 1. Collinearity Check & 2. Residual Analysis

In [20]:

```
# 1) Collinearity problem Check
# Calculate VIF = 1/(1-Rsquare) For all independent variables

rsq_age=smf.ols("Age~KM+HP+CC+Doors+Gears+QT+Weight",data=toyo4).fit().rsquared
vif_age=1/(1-rsq_age)

rsq_km=smf.ols("KM~Age+HP+CC+Doors+Gears+QT+Weight",data=toyo4).fit().rsquared
vif_km=1/(1-rsq_km)

rsq_hp=smf.ols("HP~Age+KM+CC+Doors+Gears+QT+Weight",data=toyo4).fit().rsquared
vif_hp=1/(1-rsq_hp)

rsq_cc=smf.ols("CC~Age+KM+HP+Doors+Gears+QT+Weight",data=toyo4).fit().rsquared
vif_cc=1/(1-rsq_cc)

rsq_dr=smf.ols("Doors~Age+KM+HP+CC+Gears+QT+Weight",data=toyo4).fit().rsquared
vif_dr=1/(1-rsq_dr)

rsq_gr=smf.ols("Gears~Age+KM+HP+CC+Doors+QT+Weight",data=toyo4).fit().rsquared
vif_gr=1/(1-rsq_gr)

rsq_qt=smf.ols("QT~Age+KM+HP+CC+Doors+Gears+Weight",data=toyo4).fit().rsquared
vif_qt=1/(1-rsq_qt)

rsq_wt=smf.ols("Weight~Age+KM+HP+CC+Doors+Gears+QT",data=toyo4).fit().rsquared
vif_wt=1/(1-rsq_wt)

# Putting the values in DataFrame
d1={'Vriables':['Age','KM','HP','CC','Doors','Gears','QT','Weight'],
    'VIF':[vif_age ,vif_km,vif_hp,vif_cc,vif_dr,vif_gr,vif_qt,vif_wt]}
Vif_df=pd.DataFrame(d1)
Vif_df
```

Out[20]:


	Vriables	VIF
0	Age	1.876236
1	KM	1.757178
2	HP	1.419180
3	CC	1.163470
4	Doors	1.155890
5	Gears	1.098843
6	QT	2.295375
7	Weight	2.487180

**None variable has VIF>20, No Collinearity,so consider all variables in Regression**

## equation

In [21]:

```
# 2) Residual Analysis
# Test for normality of residuals (Q-Q Plot) using residual model (model.resid)
sm.qqplot(model.resid, line='q') # 'q' - A line is fit through the quantiles # line= '45'-
plt.title("Normal Q-Q plot of residuals")
plt.show()
```



In [22]:

```
list(np.where(model.resid > 6000)) # outlier detection from above Q-Q plot of residual
```

Out[22]:

```
[array([109, 146, 522], dtype=int64)]
```

In [23]:

```
list(np.where(model.resid < -6000))
```

Out[23]:


```
[array([220, 600, 959], dtype=int64)]
```

In [24]:

```
# Test for Homoscedasticity or heteroscedasticity (plotting model's standardized fitted val
def standard_values(vals) : return (vals - vals.mean()) / vals.std() # user defined z= (x-mu)/
```

In [25]:

```
plt.scatter(standard_values(model.fittedvalues),standard_values(model.resid))  
plt.title('Residual Plot')  
plt.xlabel('standardized fitted values')  
plt.ylabel('standardized residual values')  
plt.show()
```

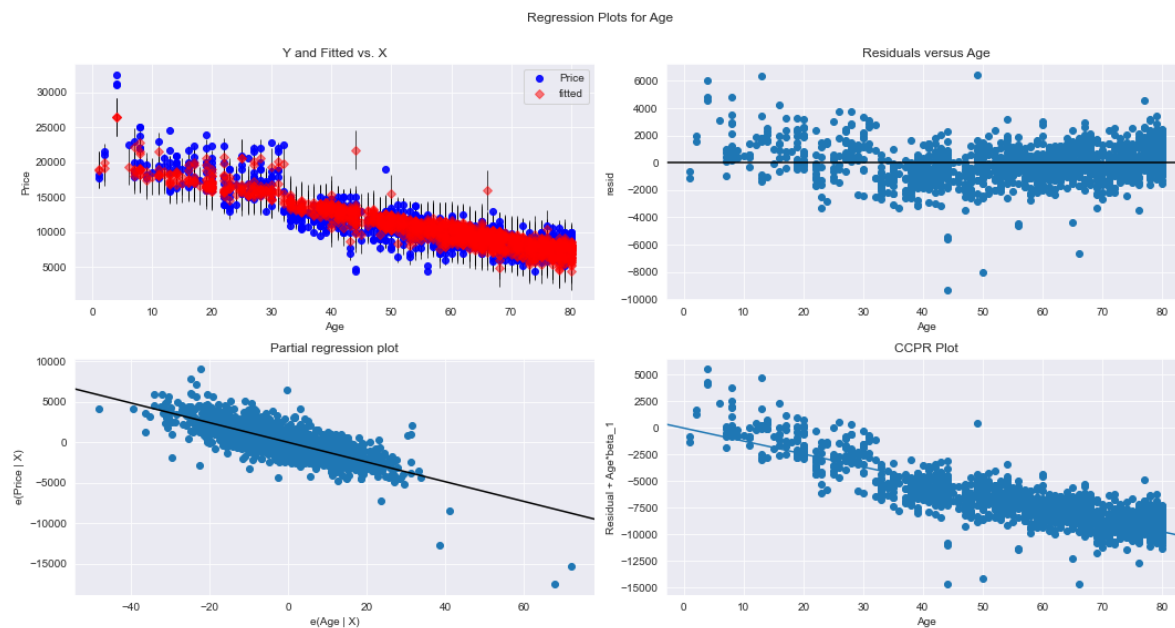


In [26]:

```
# Test for error or Residuals Vs Regressors or Independent 'x'variables or predictors  
# Using Residual Regression plot code graphics.plot_regress_exog(model,'x',fig) # exog = x-
```

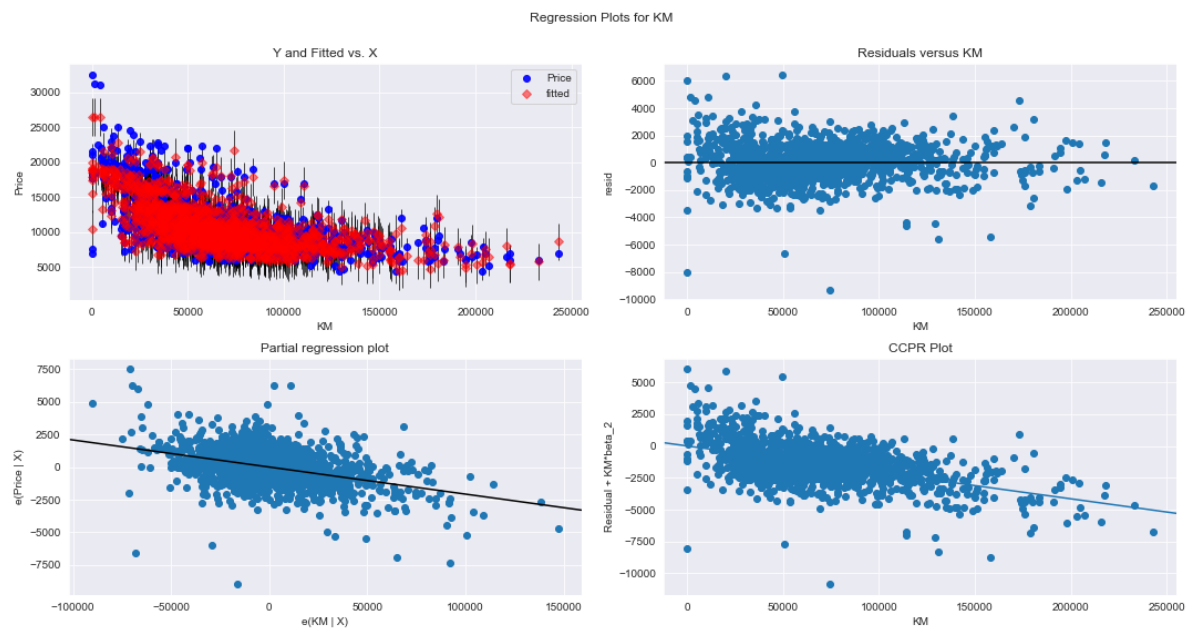
In [27]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Age',fig=fig)
plt.show()
```



In [28]:

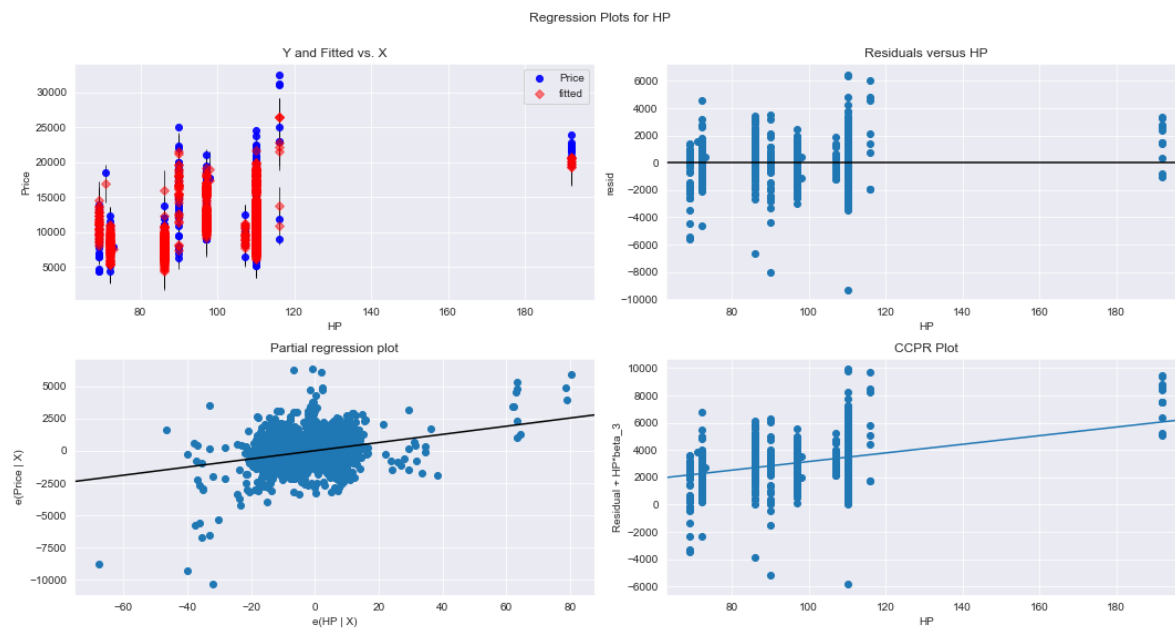
```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'KM',fig=fig)
plt.show()
```





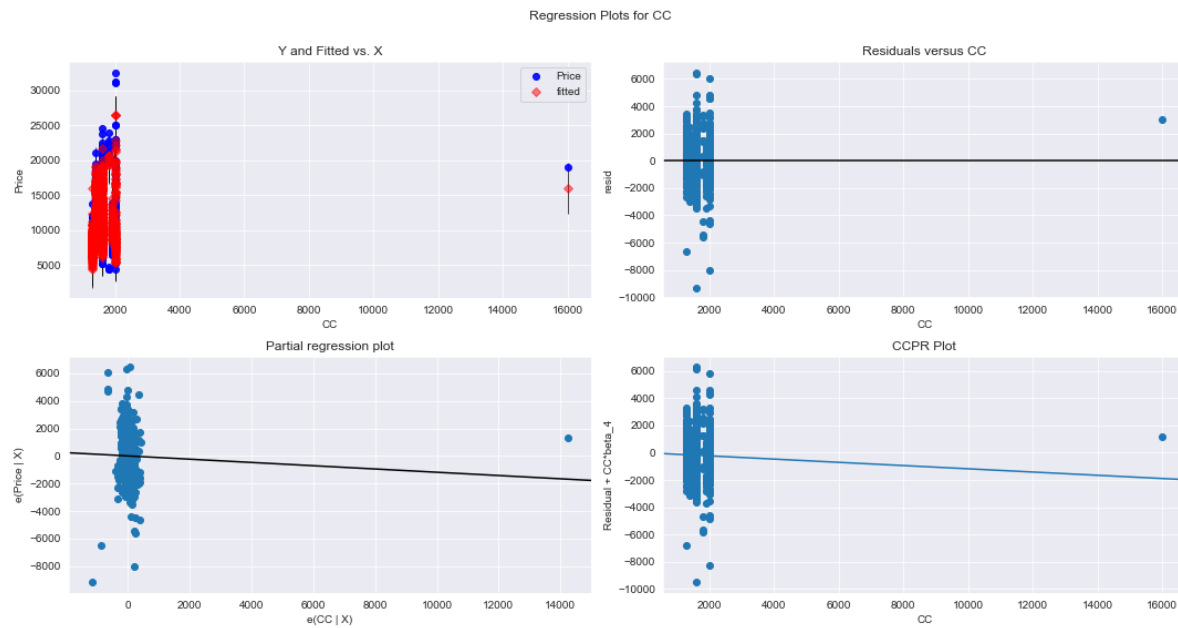
In [29]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'HP',fig=fig)
plt.show()
```



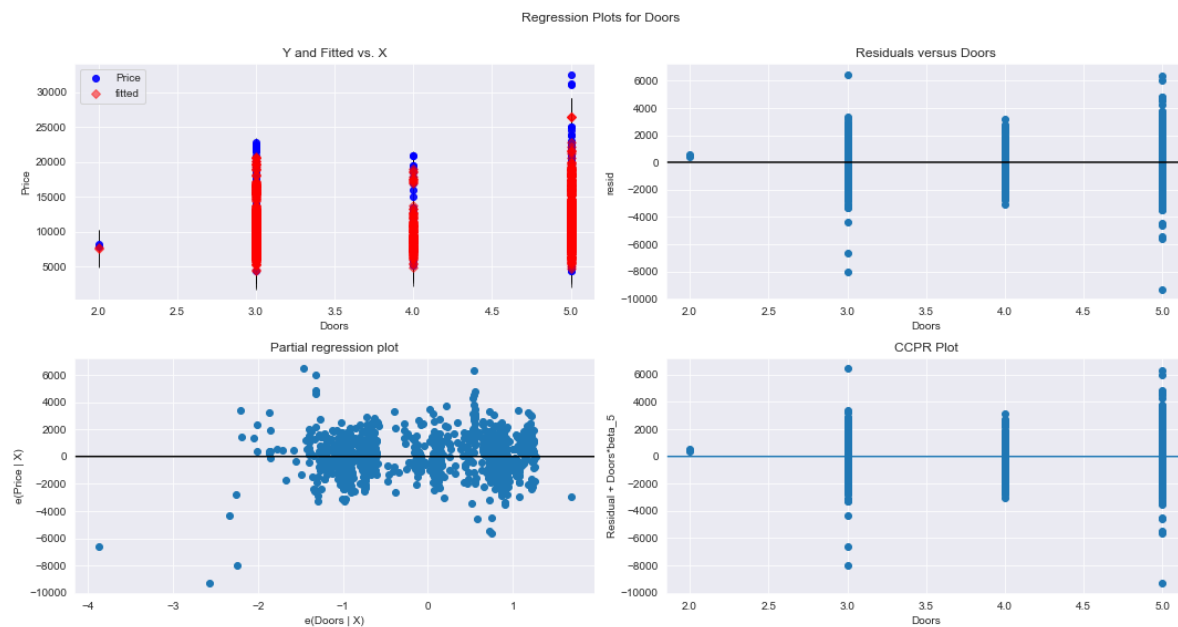
In [30]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'CC',fig=fig)
plt.show()
```



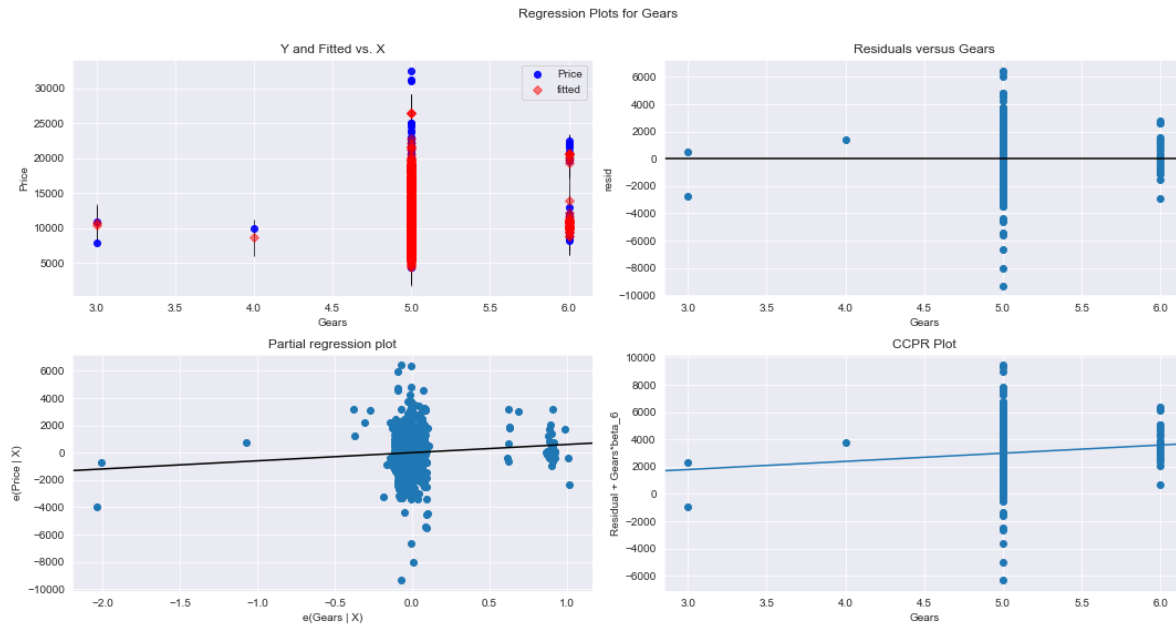
In [31]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Doors',fig=fig)
plt.show()
```



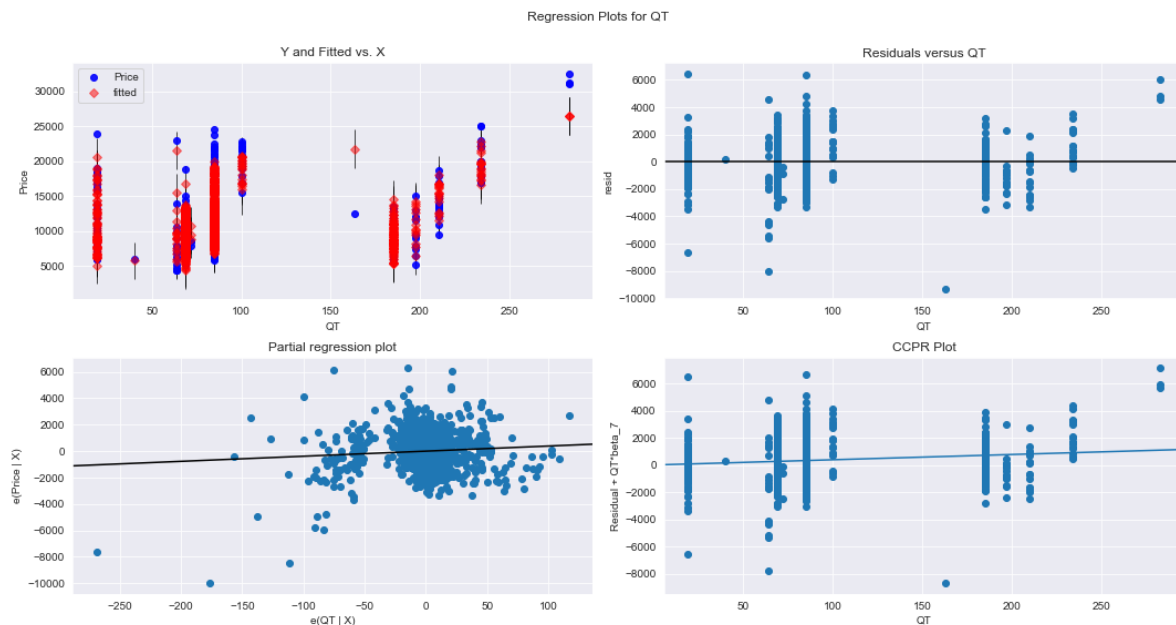
In [32]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Gears',fig=fig)
plt.show()
```



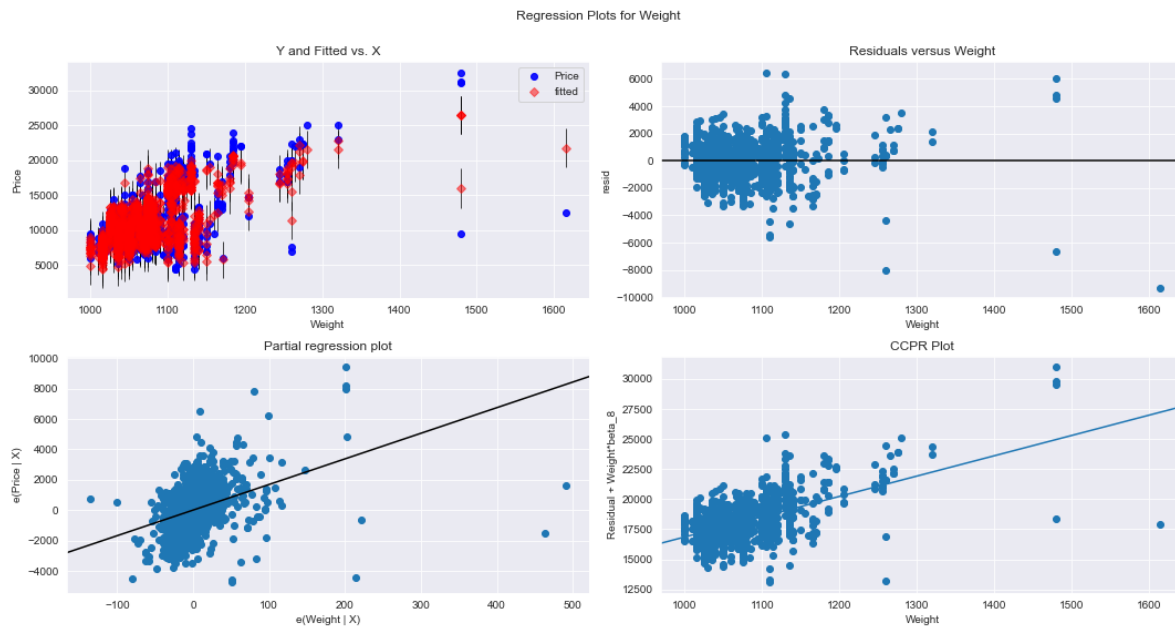
In [33]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'QT',fig=fig)
plt.show()
```



In [34]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'Weight',fig=fig)
plt.show()
```



## Model Deletion Diagnostics (checking outliers or influencers)

Two Techniques: 1.Cook's Distance & 2. Leverage value

In [35]:


```
# Cook's Distance: If Cook's distance > 1, Then it's an outlier
# Get influencers using cook's distance
(c,_)=model.get_influence().cooks_distance
c
```

Out[35]:

```
array([7.22221054e-03, 3.94547973e-03, 5.44224039e-03, ...,
       8.04110550e-07, 6.99854767e-04, 1.08408002e-02])
```

In [36]:

```
# Plot the influencer using the stem plot
fig=plt.figure(figsize=(20,7))
plt.stem(np.arange(len(toyo4)),np.round(c,3))
plt.xlabel('Row Index')
plt.ylabel('Cooks Distance')
plt.show()
```



In [37]:


```
# Index and value of influencer where C>0.5
np.argmax(c),np.max(c)
```

Out[37]:

(80, 78.7295058224916)

In [38]:

```
# 2.Leverage Value using High Influncer Points: Points beyond Leverage_cutoff value are inf
fig,ax=plt.subplots(figsize=(20,20))
fig=influence_plot(model,ax = ax)
```



In [39]:

```
# Leverage Cutoff value =  $3 \cdot (k+1) / n$ ;  $k$  = no. of features/columns &  $n$  = no. of datapoints  
k=toyo4.shape[1]  
n=toyo4.shape[0]  
leverage_cutoff=(3*(k+1))/n  
leverage_cutoff
```

Out[39]:

0.020905923344947737

In [40]:

```
toyo4[toyo4.index.isin([80])]
```

Out[40]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
80	18950	25	20019	110	16000	5	5	100	1180

## improving the model

In [41]:

```
# Creating a copy of data so that original dataset is not affected
toyo_new=toyo4.copy()
toyo_new
```

Out[41]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...	...	...	...	...	...	...	...	...	...
1430	7500	69	20544	86	1300	3	5	69	1025
1431	10845	72	19000	86	1300	3	5	69	1015
1432	8500	71	17016	86	1300	3	5	69	1015
1433	7250	70	16916	86	1300	3	5	69	1015
1434	6950	76	1	110	1600	5	5	19	1114

1435 rows × 9 columns

In [42]:

```
# Discard the data points which are influencers and reassign the row number (reset_index(drop=True))
toyo5=toyo_new.drop(toyo_new.index[[80]],axis=0).reset_index(drop=True)
toyo5
```

Out[42]:

	Price	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	13500	23	46986	90	2000	3	5	210	1165
1	13750	23	72937	90	2000	3	5	210	1165
2	13950	24	41711	90	2000	3	5	210	1165
3	14950	26	48000	90	2000	3	5	210	1165
4	13750	30	38500	90	2000	3	5	210	1170
...	...	...	...	...	...	...	...	...	...
1429	7500	69	20544	86	1300	3	5	69	1025
1430	10845	72	19000	86	1300	3	5	69	1015
1431	8500	71	17016	86	1300	3	5	69	1015
1432	7250	70	16916	86	1300	3	5	69	1015
1433	6950	76	1	110	1600	5	5	19	1114

1434 rows × 9 columns



## Model Deletion Diagnostics and Final Model

In [43]:

```
while np.max(c)>0.5:
    model=smf.ols("Price~Age+KM+HP+CC+Doors+Gears+QT+Weight",data=toyo5).fit()
    (c,_)=model.get_influence().cooks_distance
    c
    np.argmax(c),np.max(c)
    toyo5=toyo5.drop(toyo5.index[[np.argmax(c)]],axis=0).reset_index(drop=True)
    toyo5
else:
    final_model=smf.ols("Price~Age+KM+HP+CC+Doors+Gears+QT+Weight",data=toyo5).fit()
    final_model.rsquared, final_model.aic
    print("Thus model accuracy is improved to",final_model.rsquared)
```

Thus model accuracy is improved to 0.8882395145171204

In [44]:

```
if np.max(c)>0.5:
    model=smf.ols("Price~Age+KM_HP+CC+Doors+Gears+QT+Weight",data=toyo5).fit()
    (c,_)=model.get_influence().cooks_distance
    c

    np.argmax(c),np.max(c)
    toyo5=toyo5.drop(toyo5.index[[np.argmax(c)]],axis=0).reset_index(drop=True)
    toyo5
else:
    final_model=smf.ols("Price~Age+KM+HP+CC+Doors+Gears+QT+Weight",data=toyo5).fit()
    final_model.rsquared, final_model.aic
    print("Thus model accuracy is improved to",final_model.rsquared)
```

Thus model accuracy is improved to 0.8882395145171204

In [45]:

```
final_model.rsquared
```

Out[45]:

0.8882395145171204

## Creation of model using Lasso

In [46]:

```

x=toyo5[['Age', 'Weight', 'KM', 'HP', 'CC', 'Doors', 'QT', 'Weight']]
y=toyo5['Price']

from sklearn.model_selection import train_test_split
from sklearn.linear_model import Lasso
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=10)
reg=Lasso(alpha=0.5)
reg.fit(x_train, y_train)

print('Lasso Regression: R^2 score on training set',reg.score(x_train, y_train)*100)
print('Lasso Regression: R^2 score on test set',reg.score(x_test, y_test)*100)

lambdas=(0.001,0.01,0.1,0.5,1,2,10)
l_num=7
pred_num=x.shape[1]

# prepare data for enumerate
coeff_a = np.zeros((l_num, pred_num))
train_r_squared =np.zeros(l_num)
test_r_squared =np.zeros(l_num)

# enumerate through Lambdas with index and i
for ind, i in enumerate(lambdas):
    reg = Lasso(alpha =1)
    reg.fit(x_train, y_train)

    coeff_a[ind,:]=reg.coef_
    train_r_squared[ind] = reg.score(x_train, y_train)
    test_r_squared[ind] = reg.score(x_test, y_test)

```

Lasso Regression: R^2 score on training set 88.74698071262482

Lasso Regression: R^2 score on test set 88.77836439214997

## Creation of model using Ridge

In [47]:

```

from sklearn.linear_model import Ridge
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=10)
reg=Ridge(alpha=0.5)
reg.fit(x_train, y_train)

print('Ridge Regression: R^2 score on training set',reg.score(x_train, y_train)*100)
print('Ridge Regression: R^2 score on test set',reg.score(x_test, y_test)*100)

lambdas=(0.001,0.01,0.1,0.5,1,2,10)
l_num=7
pred_num=x.shape[1]

# prepare data for enumerate
coeff_a = np.zeros((l_num, pred_num))
train_r_squared =np.zeros(l_num)
test_r_squared =np.zeros(l_num)

# enumerate through Lambdas with index and i
for ind, i in enumerate(lambdas):
    reg = Ridge(alpha =1)
    reg.fit(x_train, y_train)

    coeff_a[ind,:]=reg.coef_
    train_r_squared[ind] = reg.score(x_train, y_train)
    test_r_squared[ind] = reg.score(x_test, y_test)

```

Ridge Regression: R^2 score on training set 88.74698328222092

Ridge Regression: R^2 score on test set 88.77884476539593

## Model Predictions

In [48]:

```

# say New data for prediction is
new_data=pd.DataFrame({"Age":12,"KM":40000,"HP":80,"CC":1300,"Doors":4,"Gears":5,"QT":69,"W
new_data

```

Out[48]:

	Age	KM	HP	CC	Doors	Gears	QT	Weight
0	12	40000	80	1300	4	5	69	1012

In [49]:

```

# Manual prediction of price
final_model.predict(new_data)

```

Out[49]:

```

0    14341.570181
dtype: float64

```

In [50]:

```
# Autotamatioc precdiction of price with 90.02% accuracy
pred_y=final_model.predict(toyo5)
pred_y
```

Out[50]:

```
0      16345.352610
1      15886.635544
2      16328.224968
3      15996.318854
4      15883.424182
...
1426    9161.230587
1427    8536.091326
1428    8681.531063
1429    8793.668694
1430   10860.695492
Length: 1431, dtype: float64
```

## Table Containig R^2 value for each prepared model

In [51]:

```
d2={'Prep_Models':['Model','Final_Model'],'Rsquared':[model.rsquared,final_model.rsquared]}
table=pd.DataFrame(d2)
table
```

Out[51]:

	Prep_Models	Rsquared
0	Model	0.883968
1	Final_Model	0.888240

In [ ]:

## Assignment-05-Multiple Linear Regression-2

In [52]:

```
#Import dataset
data=pd.read_csv("C:/Users/LENOVO/Documents/Custom Office Templates/50_Startups.csv")
data
```

Out[52]:

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84

	R&D Spend	Administration	Marketing Spend	State	Profit
33	55493.95	103057.49	214634.81	Florida	96778.92
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

In [53]:

# EDA

data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   R&D Spend              50 non-null     float64
1   Administration         50 non-null     float64
2   Marketing Spend        50 non-null     float64
3   State                  50 non-null     object
4   Profit                 50 non-null     float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

In [54]:

data.isnull().sum()

Out[54]:

```
R&D Spend      0
Administration 0
Marketing Spend 0
State          0
Profit         0
dtype: int64
```

Inference:

No NA values

In [55]:

```
data1=data.rename({'R&D Spend':'RDS','Administration':'ADMS','Marketing Spend':'MKTS'},axis=1)
data1
```

Out[55]:

	RDS	ADMS	MKTS	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92



	RDS	ADMS	MKTS	State	Profit
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

In [56]:

```
data1[data1.duplicated()] # No duplicated data
```

Out[56]:

RDS	ADMS	MKTS	State	Profit
-----	------	------	-------	--------

In [57]:

```
data1.describe()
```

Out[57]:

	RDS	ADMS	MKTS	Profit
count	50.000000	50.000000	50.000000	50.000000
mean	73721.615600	121344.639600	211025.097800	112012.639200
std	45902.256482	28017.802755	122290.310726	40306.180338
min	0.000000	51283.140000	0.000000	14681.400000
25%	39936.370000	103730.875000	129300.132500	90138.902500
50%	73051.080000	122699.795000	212716.240000	107978.190000
75%	101602.800000	144842.180000	299469.085000	139765.977500
max	165349.200000	182645.560000	471784.100000	192261.830000

# Correlation Analysis

In [58]:

data1.corr()

Out[58]:

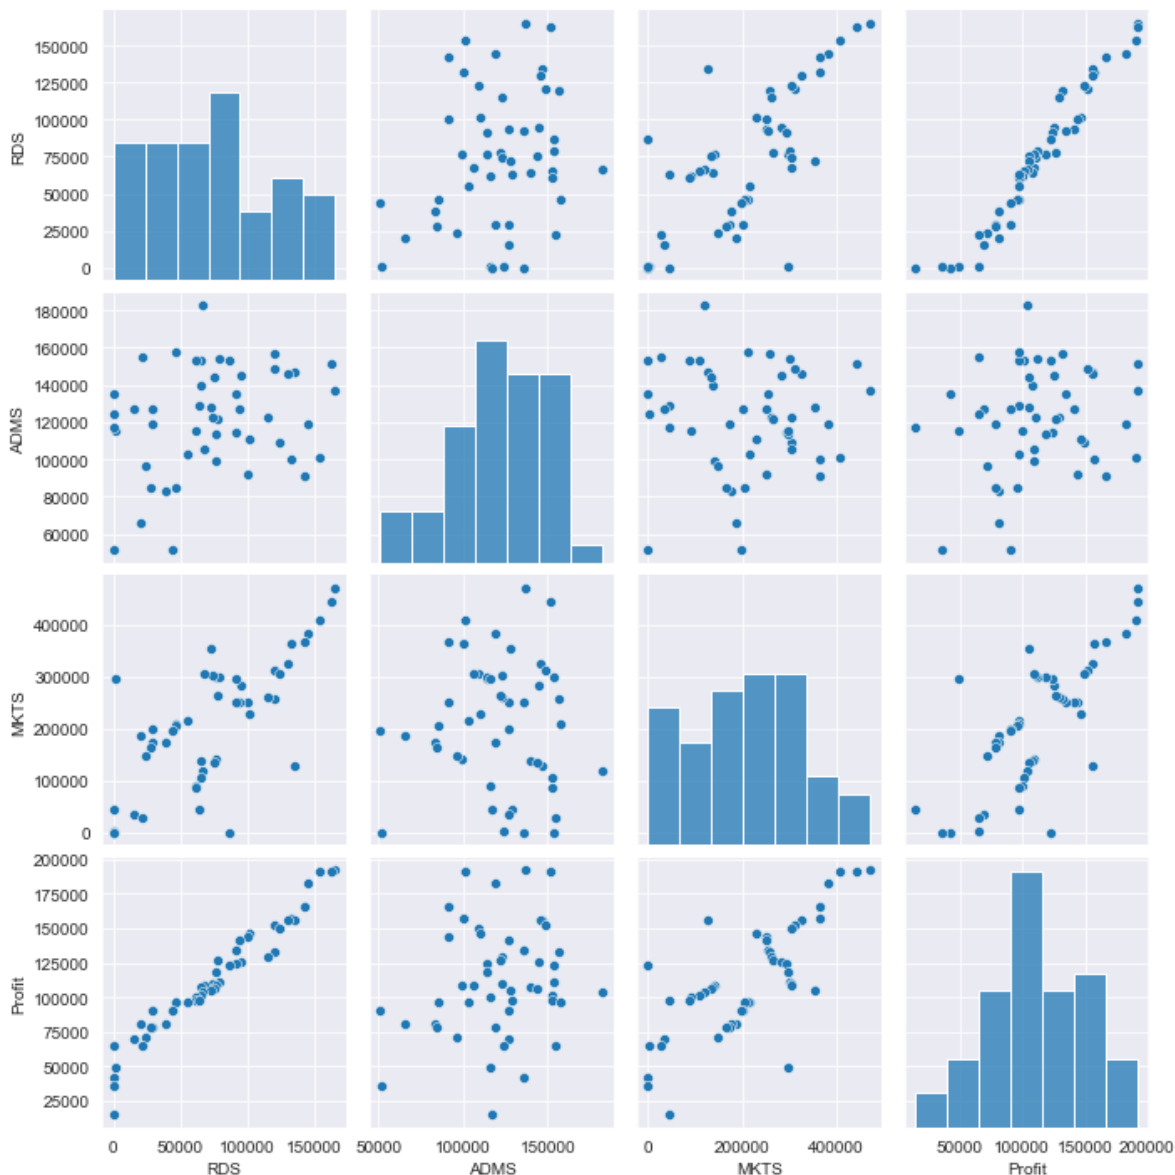
	RDS	ADMS	MKTS	Profit
RDS	1.000000	0.241955	0.724248	0.972900
ADMS	0.241955	1.000000	-0.032154	0.200717
MKTS	0.724248	-0.032154	1.000000	0.747766
Profit	0.972900	0.200717	0.747766	1.000000

In [59]:

```
sns.set_style(style='darkgrid')
sns.pairplot(data1)
```

Out[59]:

&lt;seaborn.axisgrid.PairGrid at 0x74e7940310&gt;



# Model Building

In [60]:

```
model=smf.ols("Profit~RDS+ADMS+MKTS",data=data1).fit()
```

In [61]:

```
# Model Testing  
# finding coefficient parameters  
model.params
```

Out[61]:

```
Intercept    50122.192990  
RDS           0.805715  
ADMS         -0.026816  
MKTS          0.027228  
dtype: float64
```

In [62]:

```
# finding tvalues and pvalues  
model.tvalues, np.round(model.pvalues,5)
```

Out[62]:

```
(Intercept    7.626218  
RDS          17.846374  
ADMS         -0.525507  
MKTS          1.655077  
dtype: float64,  
Intercept    0.00000  
RDS           0.00000  
ADMS          0.60176  
MKTS          0.10472  
dtype: float64)
```

Inference:

ADMS and MKTS has insignificant values

In [63]:

```
# Finding rsquared values  
model.rsquared, model.rsquared_adj # Model accuracy is 94.75%
```

Out[63]:

```
(0.9507459940683246, 0.9475337762901719)
```

In [64]:

```
# Build SLR and MLR models for insignificant variables 'ADMS' and 'MKTS'  
# Also find their tvalues and pvalues
```

In [65]:

```
slr_a=smf.ols("Profit~ADMS", data=data1).fit()  
slr_a.tvalues, slr_a.pvalues # ADMS has insignificant pvalue
```

Out[65]:

```
(Intercept    3.040044  
ADMS          1.419493  
dtype: float64,  
Intercept    0.003824  
ADMS          0.162217  
dtype: float64)
```

In [66]:

```
slr_m=smf.ols("Profit~MKTS", data=data1).fit()  
slr_m.tvalues, slr_m.pvalues #MKTS has significant pvalue
```

Out[66]:

```
(Intercept    7.808356  
MKTS          7.802657  
dtype: float64,  
Intercept    4.294735e-10  
MKTS          4.381073e-10  
dtype: float64)
```

In [67]:

```
mlr_am=smf.ols("Profit~ADMS+MKTS", data=data1).fit()  
mlr_am.tvalues, mlr_am.pvalues # variables have significant pvalues
```

Out[67]:

```
(Intercept    1.142741  
ADMS          2.467779  
MKTS          8.281039  
dtype: float64,  
Intercept    2.589341e-01  
ADMS          1.729198e-02  
MKTS          9.727245e-11  
dtype: float64)
```

## Model Validation

### Two Techniques: 1. Collinearity Check & 2. Residual Analysis

In [68]:

```

# 1) Collinearity Problem Check
# Calculate VIF =1/(1-Rsquare) for all independent variables

rsq_r=smf.ols("RDS~ADMS+MKTS",data=data1).fit().rsquared
vif_r=1/(1-rsq_r)

rsq_a=smf.ols("ADMS~RDS+MKTS",data=data1).fit().rsquared
vif_a=1/(1-rsq_a)

rsq_m=smf.ols("MKTS~RDS+ADMS",data=data1).fit().rsquared
vif_m=1/(1-rsq_m)

# Putting the values in DataFrame
d1={'Variables':['RDS','ADM','MKTS'], 'VIF':[vif_r,vif_a,vif_m]}
Vif_df=pd.DataFrame(d1)
Vif_df

```

Out[68]:

	Variables	VIF
0	RDS	2.468903
1	ADM	1.175091
2	MKTS	2.326773

In [69]:

```

# None variable has Vif>20, No Collinearity, so consider all variables in Regression equation

```

In [70]:

```
# 2) Residual analysis
# Test for Normality of Residual (Q-Q Plot) using residual model(model.resid)

sm.qqplot(model.resid,line='q')
plt.title('Normal Q-Q plot of Residual')
plt.show()
```



In [71]:

```
list(np.where(model.resid<-30000))
```

Out[71]:


```
[array([49], dtype=int64)]
```

In [72]:

```
# Test for Homoscedasticity or heteroscedasticity (plotting model's standardized fitted val
def standard_values(vals) : return (vals-vals.mean())/vals.std() # user defined z= (x-mu)/
```

In [73]:

```
plt.scatter(standard_values(model.fittedvalues),standard_values(model.resid))
plt.title('Residual Plot')
plt.xlabel('standardized fitted values')
plt.ylabel('standardized residual values')
plt.show()
```

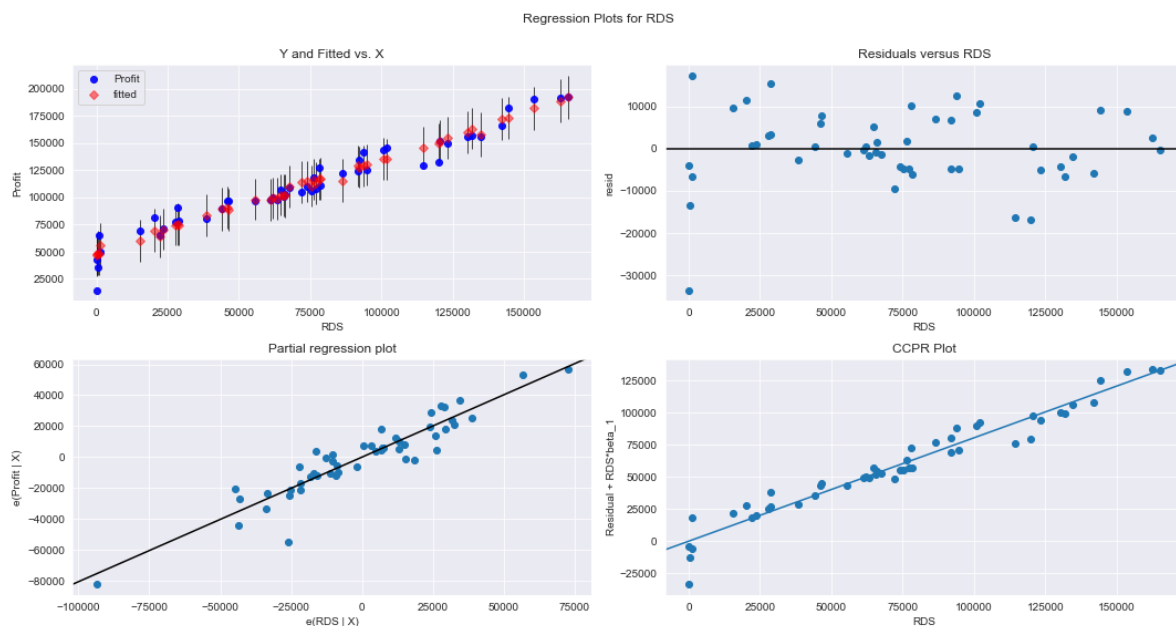


In [74]:

```
# Test for error or Residuals Vs Regressors or Independent 'x' variables or predictors
# Using Residual Regression plot code graphics.plot_regress_exog(model, 'x',fig) # exog = x-
```

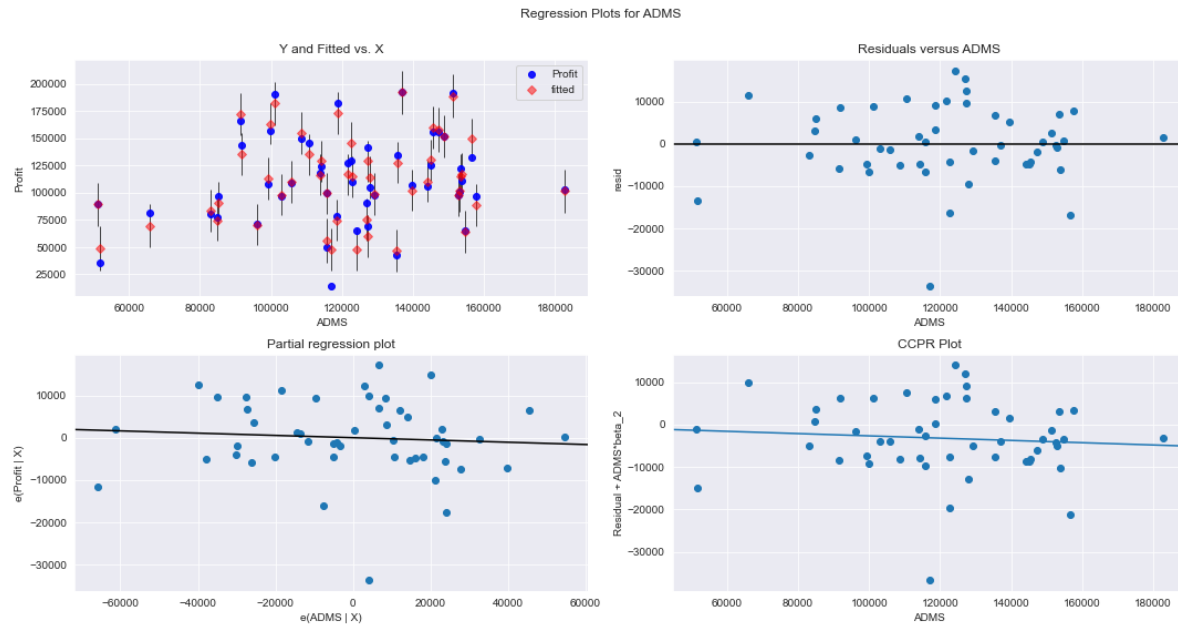
In [75]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'RDS',fig=fig)
plt.show()
```



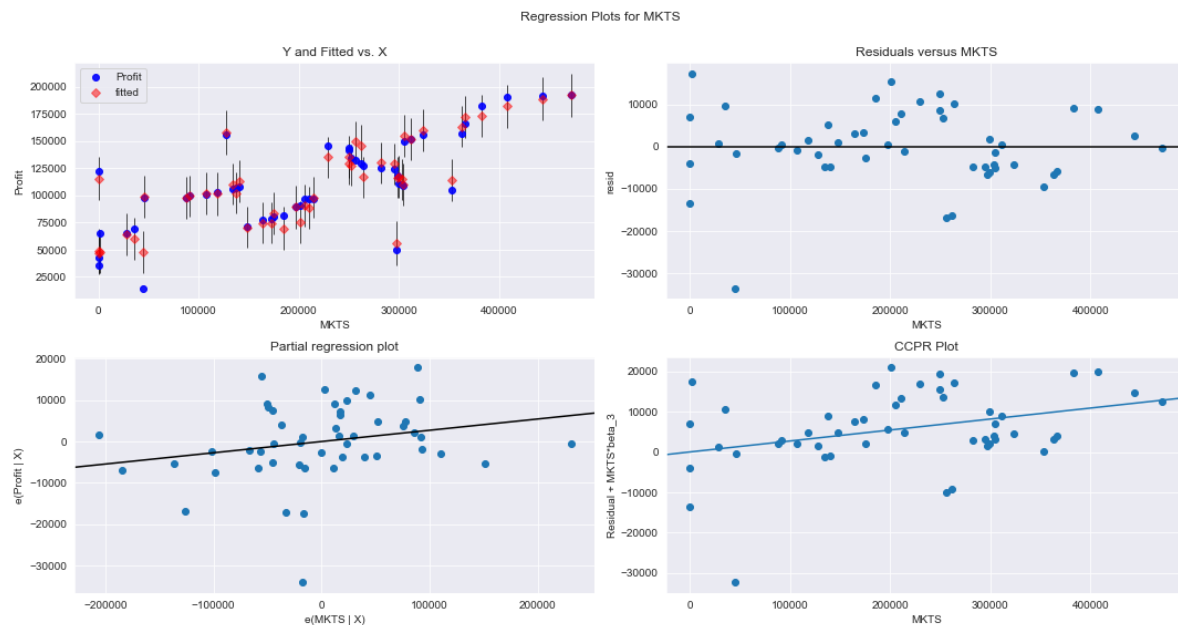
In [76]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'ADMS', fig=fig)
plt.show()
```



In [77]:

```
fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model, 'MKTS', fig=fig)
plt.show()
```



## Model Deletion Diagnostics (checking outliers or influencers)



## Two Techniques: 1.Cook's Distance & 2. Leverage value

In [78]:


```
# Cook's Distance: If Cook's distance > 1, Then it's an outlier
# Get influencers using cook's distance
(c,_)=model.get_influence().cooks_distance
c
```

Out[78]:

```
array([3.21825244e-05, 3.27591036e-03, 3.23842699e-02, 2.17206555e-02,
       1.44833032e-02, 1.17158463e-02, 2.91766303e-03, 3.56513444e-03,
       4.04303948e-05, 4.86758017e-03, 1.51064757e-02, 1.63564959e-02,
       1.15516625e-02, 4.01422811e-03, 6.12934253e-02, 3.40013448e-02,
       8.33556413e-03, 3.30534399e-03, 2.16819303e-03, 4.07440577e-02,
       4.25137222e-04, 1.09844352e-02, 2.91768000e-03, 2.76030254e-04,
       5.04643588e-03, 3.00074623e-03, 3.41957068e-03, 2.98396413e-02,
       1.31590664e-03, 1.25992620e-04, 4.18505125e-05, 9.27434786e-06,
       7.08656521e-04, 1.28122674e-04, 2.09815032e-02, 6.69508674e-03,
       5.55314705e-02, 6.55050578e-05, 5.61547311e-02, 1.54279607e-03,
       1.84850929e-03, 1.97578066e-03, 1.36089280e-04, 2.05553171e-02,
       1.23156041e-04, 9.03234206e-02, 5.45303387e-02, 5.33885616e-03,
       1.90527441e-01, 2.88082293e-01])
```

In [79]:

```
# Plot the influencer using the stem plot
fig=plt.figure(figsize=(20,7))
plt.stem(np.arange(len(data1)),np.round(c,5))
plt.xlabel('Row Index')
plt.ylabel('Cooks Distance')
plt.show()
```



In [80]:


```
# Index and value of influncer where C>0.5
np.argmax(c),np.max(c)
```

Out[80]:

```
(49, 0.28808229275432634)
```

In [81]:

```
# 2. Leverage Value using High Influence Points : Points beyond Leverage_cutoff value are i
influence_plot(model)
plt.show()
```



In [82]:

```
# Leverage Cutoff value = 3*(k+1)/n; k= no.of features/columns & n=no.of datapoints
k=data1.shape[1]
n=data1.shape[0]
leverage_cutoff=(3*(k+1))/n
leverage_cutoff
```

Out[82]:

0.36

In [83]:

```
data1[data1.index.isin([49])]
```

Out[83]:

	RDS	ADMS	MKTS	State	Profit
49	0.0	116983.8	45173.06	California	14681.4

In [84]:

```
# Discard the data points which are influencers and reassign the row number (reset_index(drop=True))
data2=data1.drop(toyo_new.index[[49]],axis=0).reset_index(drop=True)
data2
```

Out[84]:

	RDS	ADMS	MKTS	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84

	RDS	ADMS	MKTS	State	Profit
33	55493.95	103057.49	214634.81	Florida	96778.92
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41

## Model Deletion Diagnostics and Final Model

In [85]:

```
while np.max(c)>0.5:
    model=smf.ols("Profit~RDS+ADMS+MKTS",data=data2).fit()
    (c,_)=model.get_influence().cooks_distance
    c
    np.argmax(c),np.max(c)
    data2=data2.drop(data2.index[[np.argmax(c)]],axis=0).reset_index(drop=True)
    data2
else:
    final_model=smf.ols("Profit~RDS+ADMS+MKTS",data=data2).fit()
    final_model.rsquared, final_model.aic
    print("Thus model accuracy is improved to",final_model.rsquared)
```

Thus model accuracy is improved to 0.9613162435129847

In [86]:

```
final_model.rsquared
```

Out[86]:

0.9613162435129847

## Creation of model using Lasso

In [87]:

```
x=data2[['RDS','ADMS','MKTS']]
y=data2['Profit']

from sklearn.model_selection import train_test_split
from sklearn.linear_model import Lasso
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=10)
reg=Lasso(alpha=0.5)
reg.fit(x_train, y_train)

print('Lasso Regression: R^2 score on training set',reg.score(x_train, y_train)*100)
print('Lasso Regression: R^2 score on test set',reg.score(x_test, y_test)*100)

lambdas=(0.001,0.01,0.1,0.5,1,2,10)
l_num=7
pred_num=x.shape[1]

# prepare data for enumerate
coeff_a = np.zeros((l_num, pred_num))
train_r_squared =np.zeros(l_num)
test_r_squared =np.zeros(l_num)

# enumerate through Lambdas with index and i
for ind, i in enumerate(lambdas):
    reg = Lasso(alpha =1)
    reg.fit(x_train, y_train)

    coeff_a[ind,:]=reg.coef_
    train_r_squared[ind] = reg.score(x_train, y_train)
    test_r_squared[ind] = reg.score(x_test, y_test)
```

Lasso Regression: R^2 score on training set 95.24694195134217

Lasso Regression: R^2 score on test set 97.33570553521325

## Creation of model using Ridge

In [88]:

```
from sklearn.linear_model import Ridge
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=10)
reg=Ridge(alpha=0.5)
reg.fit(x_train, y_train)

print('Ridge Regression: R^2 score on training set',reg.score(x_train, y_train)*100)
print('Ridge Regression: R^2 score on test set',reg.score(x_test, y_test)*100)

lambdas=(0.001,0.01,0.1,0.5,1,2,10)
l_num=7
pred_num=x.shape[1]

# prepare data for enumerate
coeff_a = np.zeros((l_num, pred_num))
train_r_squared =np.zeros(l_num)
test_r_squared =np.zeros(l_num)

# enumerate through Lambdas with index and i
for ind, i in enumerate(lambdas):
    reg = Ridge(alpha =1)
    reg.fit(x_train, y_train)

    coeff_a[ind,:]=reg.coef_
    train_r_squared[ind] = reg.score(x_train, y_train)
    test_r_squared[ind] = reg.score(x_test, y_test)
```

```
Ridge Regression: R^2 score on training set 95.24694195134217
Ridge Regression: R^2 score on test set 97.33570551715988
```

In [89]:

data2

Out[89]:

	RDS	ADMS	MKTS	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92

	RDS	ADMS	MKTS	State	Profit
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41

## Model Predictions

In [90]:

```
# say new data for prediction is
new_data=pd.DataFrame({'RDS':70000, 'ADMS':90000, 'MKTS':140000}, index=[0])
new_data
```

Out[90]:

	RDS	ADMS	MKTS
0	70000	90000	140000

In [91]:

```
# Manual Prediction of price
final_model.predict(new_data)
```

Out[91]:

```
0    108727.154753
dtype: float64
```



In [92]:

```
# Automatic prediction of price with 09.02% accuracy
pred_y=final_model.predict(data2)
pred_y
```

Out[92]:

```
0    190716.676999
1    187537.122227
2    180575.526396
3    172461.144642
4    170863.486721
5    162582.583177
6    157741.338633
7    159347.735318
8    151328.826941
9    154236.846778
10   135507.792682
11   135472.855621
12   129355.599449
13   127780.129139
14   149295.404796
15   145937.941975
16   117437.627921
17   130408.626295
18   129129.234457
19   116641.003121
20   117097.731866
21   117911.019038
22   115248.217796
23   110603.139045
24   114051.073877
25   103398.054385
26   111547.638935
27   114916.165026
28   103027.229434
29   103057.621761
30   100656.410227
31    99088.213693
32   100325.741335
33    98962.303136
34    90552.307809
35    91709.288672
36    77080.554255
37    90722.503244
38    71433.021956
39    85147.375646
40    76625.510303
41    76492.145175
42    72492.394974
43    62592.049718
44    67025.731107
45    50457.297206
46    58338.443625
47    49375.776655
48    51658.096812
dtype: float64
```

# Table containig R^2 value for each prepared model

In [93]:

```
d2={'Prep_Models':['Model','Final_Model'],'Rsquared':[model.rsquared,final_model.rsquared]}
table=pd.DataFrame(d2)
table
```

Out[93]:

	Prep_Models	Rsquared
0	Model	0.950746
1	Final_Model	0.961316

In [ ]: